

Scikit-learn Introduction

Introduction
Supervised Machine Learning
Clustering
Recommendations
Classifications

Lesson Objectives

- ◆ Understand some of the basics of Machine Learning (ML)
- ◆ Learn how Python Scikit-learn supports ML
- ◆ Become familiar with some of the algorithms in Scikit-learn

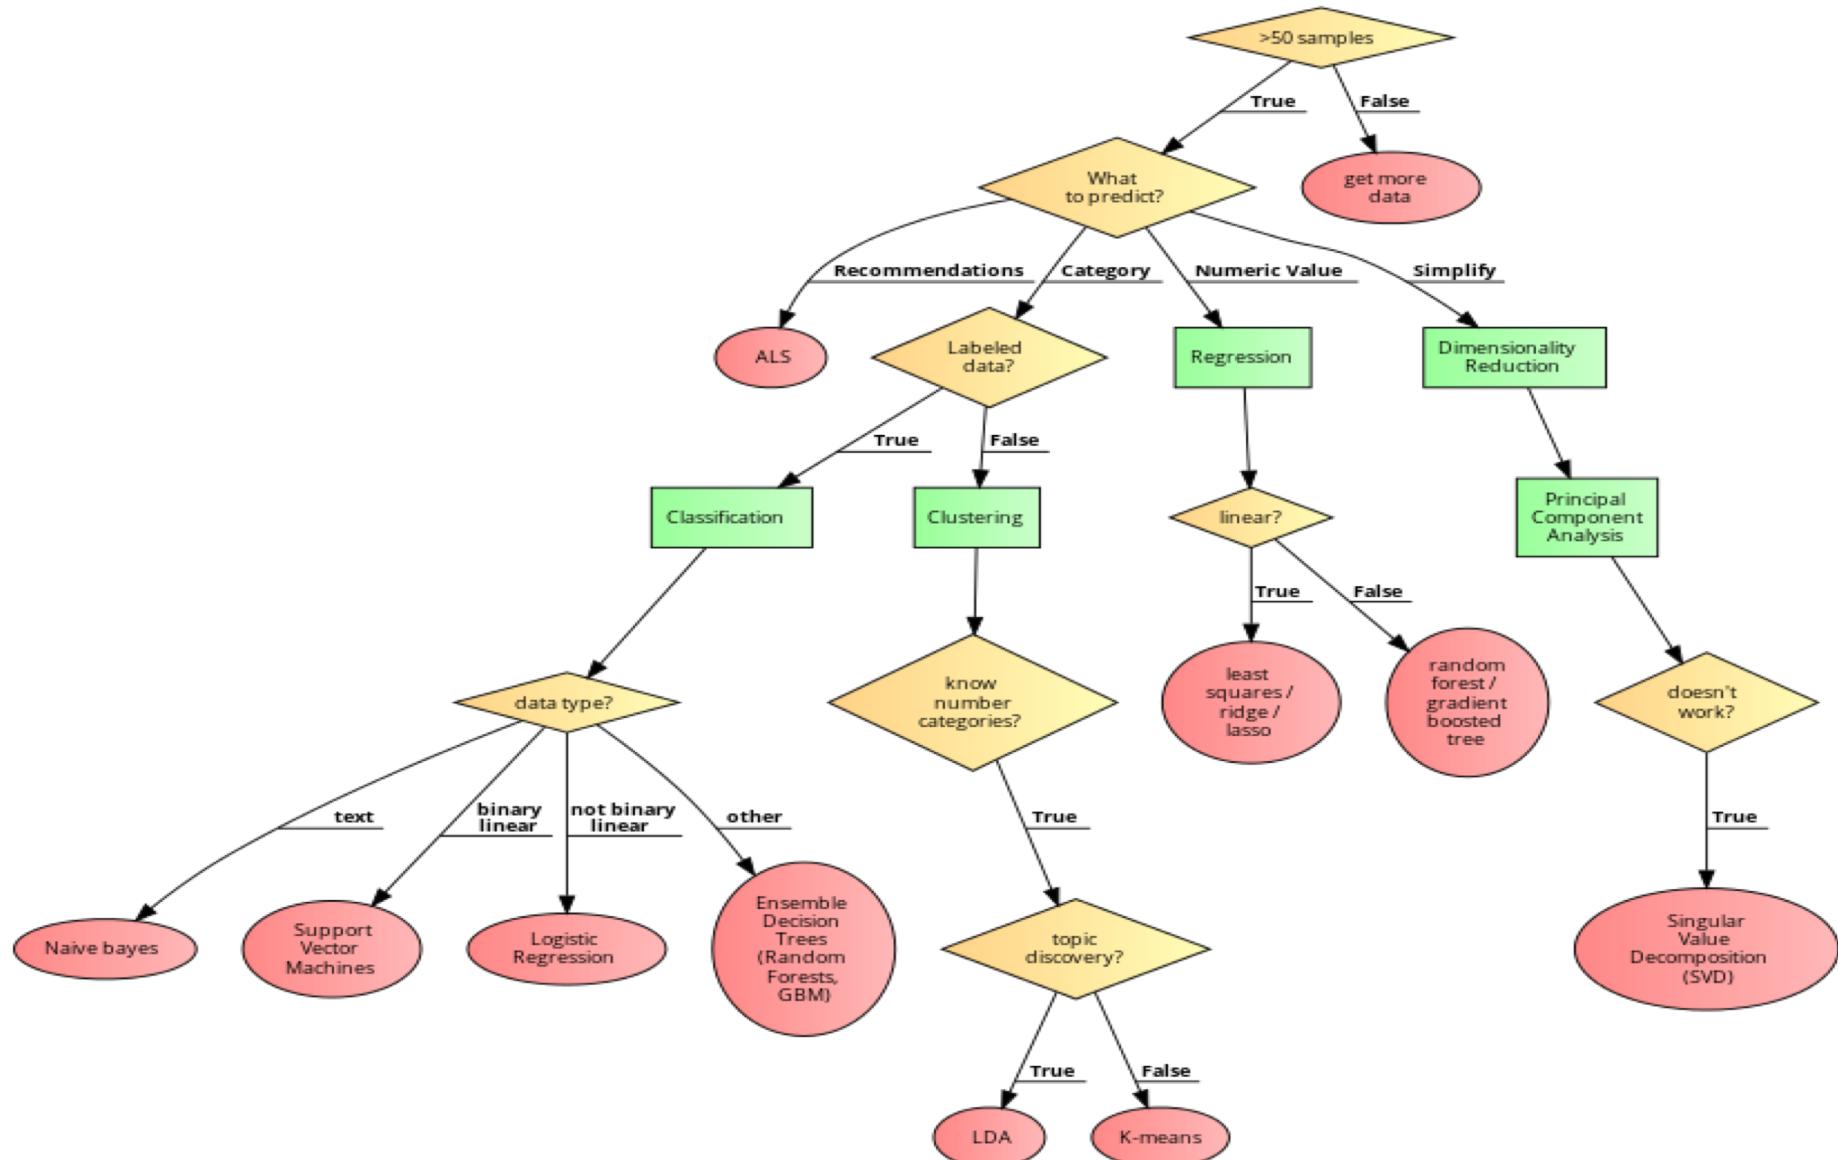
Introduction

Introduction
Supervised Machine Learning
Clustering
Recommendations
Classifications

Machine Learning in Python

- ◆ Python Scikit-learn library provides lots algorithms
- ◆ Lots of common algorithms are supported
 - Classification/Regressions
 - Linear models (linear R, logistic regression, SVM)
 - Decision trees
 - K-Means clustering
 - More

Scikit-learn Algorithm overview



Estimators (base class)

- ◆ Estimators are classes that have a `.fit()` method
- ◆ Mostly, these are classes that are used for prediction:
 - Regression
 - Classification

```
>>> estimator.fit(data)
```

↳ Estimators usually take parameters

```
>>> estimator = Estimator(param1=1,  
param2=2)  
>>> estimator.param1
```

Calling fit()

- ◆ fit can take one or two arguments
- ◆ Unsupervised learning:
 - Typically just the data
- ◆ Supervised learning: 2 params
 - Data (as ndarray)
 - Labels (the correct answer)

Transformer (base class)

- ◆ Transformers are also Estimators
 - They have a .fit() function
 - It may not necessarily be used.
- ◆ They also can transform data
 - They also have a .transform() function
- ◆ A trivial example (taking the log of the input)

```
>>> transformer =  
FunctionTransformer(np.log1p)  
>>> transformer.transform(np.array[1,2])  
[0.69314718, 1.09861229]
```

- ◆ Typically transformers are used to process our data
 - For example, to scale or normalize

Pipelines

Licensed for personal use only for Vincent Chang <vincent.chang@macys.com> from Python ML @ Macy's 04/25/2108 @

2018-04-25

- ◆ Usually we have many estimators and transformers together
- ◆ We can create a pipeline class to hold a pipeline of estimators and transformers.
- ◆ Pipeline class
- ◆ We can combine transformers together
- ◆ Or provide transformers + final estimator

Prediction

- ◆ Estimators can also have a predict() method
- ◆ Used for predicting results of a model
- ◆ An example classification model using a SVM:

```
>>> from sklearn import svm  
>>> clf = svm.SVC(gamma=0.001, C=100.)  
>>> clf.fit(data, label)  
>>> clf.predict(np.array([1,2]))  
array([3, 4]) # Prediction output
```

Feature Vectors

- ◆ Sklearn runs on Feature Vectors stored as ndArrays
- ◆ Pandas series can be converted to NdArrays if they are numeric (usually float64)
- ◆ We need to extract columns from Pandas dataframe,
 - Convert them to ndarrays
- ◆ But how do we do that?

Categorical Variables

- ◆ Algorithms run on ndarrays and matrices of type float64
 - Bool, int64 will be upcasted to float64
- ◆ Your data may have string or categorical variables
 - What to do?
- ◆ For a label:
 - Use the LabelEncoder
 - Will convert label into integer value and back.
- ◆ For a column:
 - Have to convert into a numeric

Strategy 1: Factorization

- ◆ Factorization
 - Call pd.factorize() on the dimension
 - Will assign an integer to each unique value
 - E.g. (red = 1, green = 2... etc)
- ◆ Problems with Factorization
 - Number may not be meaningful in any way (arbitrary)
 - Model may try to infer something about ordering.
 - Model may not understand that these are discrete values
 - Model may infer some correlation between values.
 - “Finding signal in the noise” – can lead to bad results

Strategy 2: Dummy Variables

- ◆ We can use Dummy Variables
 - Aka One Hot Encoding
 - This will generate One hot encoded features.

```
from sklearn.preprocessing import  
OneHotEncoder  
  
fact = pd.factorize(dimensions)  
OneHotEncoder().fit(fact).transform(fact)
```

Strategy 3: Quantization

- ◆ Turn the category into a real-world number
- ◆ Example colors:
 - Red = 450MHz
 - Green ...
- ◆ Usually join (pd.merge) to a "lookup table"

Scaling and Normalization

- ◆ Usually data needs to be cleaned up and transformed before creating features
- ◆ Scaling
 - StandardScaler class will scale data down
 - Remove high magnitude data
- ◆ Normalization
 - Will use Z-scoring to norm

Data Whitening (PCA)

Introduction
Supervised Machine Learning
Data Whitening (PCA)
Clustering
Recommendations
Classifications

Data Whitening

- ◆ Data that is heavily correlated may need to be whitened.
- ◆ This means de-correlated.
- ◆ One way to do this is with PCA: Principal Component Analysis.
- ◆ PCA can also be used to reduce dimensionality

Principal Component Analysis (PCA)

- ◆ Not all dimensions are equally interesting.
Can we reduce the '**dimensionality**' of data, **without loosing too much information?**
- ◆ PCA does just that
- ◆ It finds low-dimensional representation of data that contains as much as possible of **variation**
- ◆ PCA seeks a small number of dimensions that are as **interesting** as possible
- ◆ **Interesting** is measured by the amount that the observations vary along each dimension

Using PCA for Visualization

- ◆ If we reduce dimensions to 2 or 3, we can visualize highly dimensional data
- ◆ Great alternative to Scatter Plot Matrices

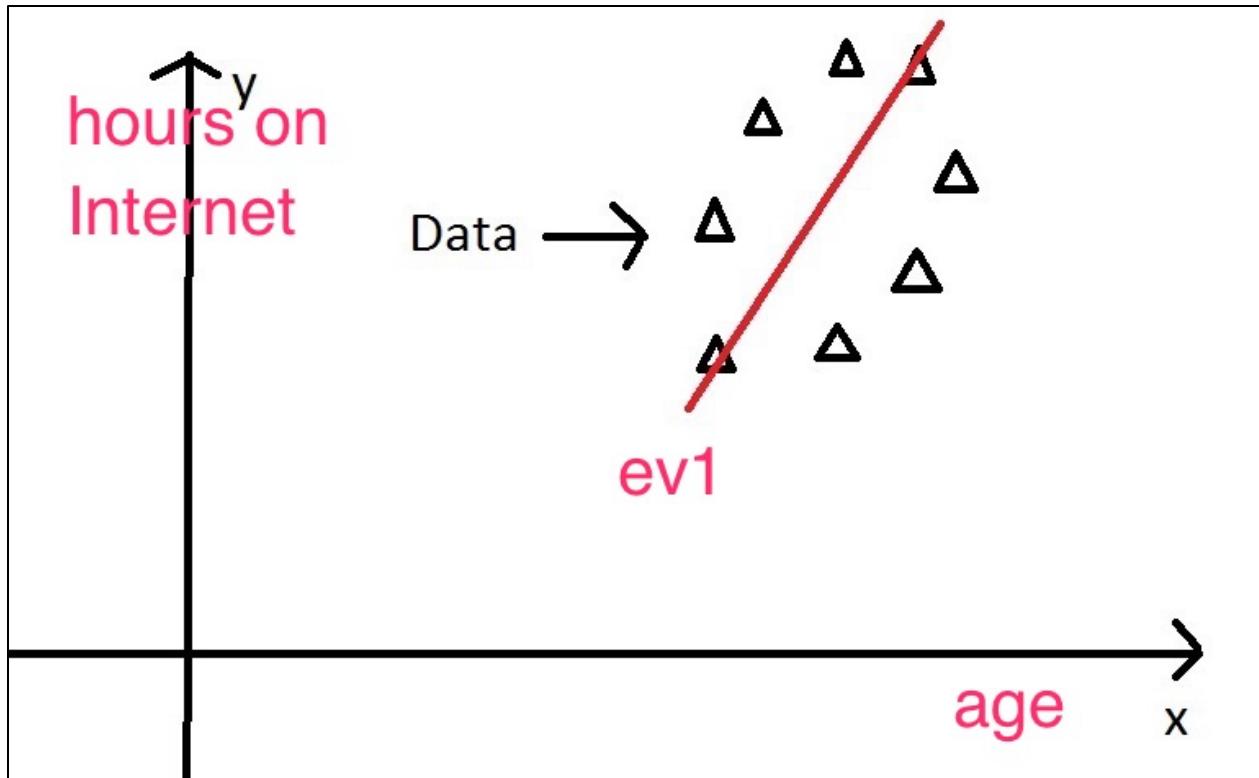
PCA Theory: Eigen Vectors / Value

- ◆ Say we are measuring the following responses
 - Age, hours on the internet
 - 2 variables → 2 dimensional dataset → 2 Eigen Vectors
- ◆ If we measure
 - Age, hours on the internet, hours on the mobile phone
 - 3 variables, 3 dimensional dataset → 3 Eigen vectors
- ◆ So number of Eigen vectors = number of dimensions

- ◆ EigenVector is a direction – vertical, horizontal, 45' degrees ..etc
- ◆ EigenValue is a number – denoting how much 'variance' in the data in that vector's direction
- ◆ Eigen Vector and Eigen Value go together
(E Vector, E Value)
- ◆ EigenVector with highest EigenValue (meaning lot of variance in that direction) becomes a Principal Component

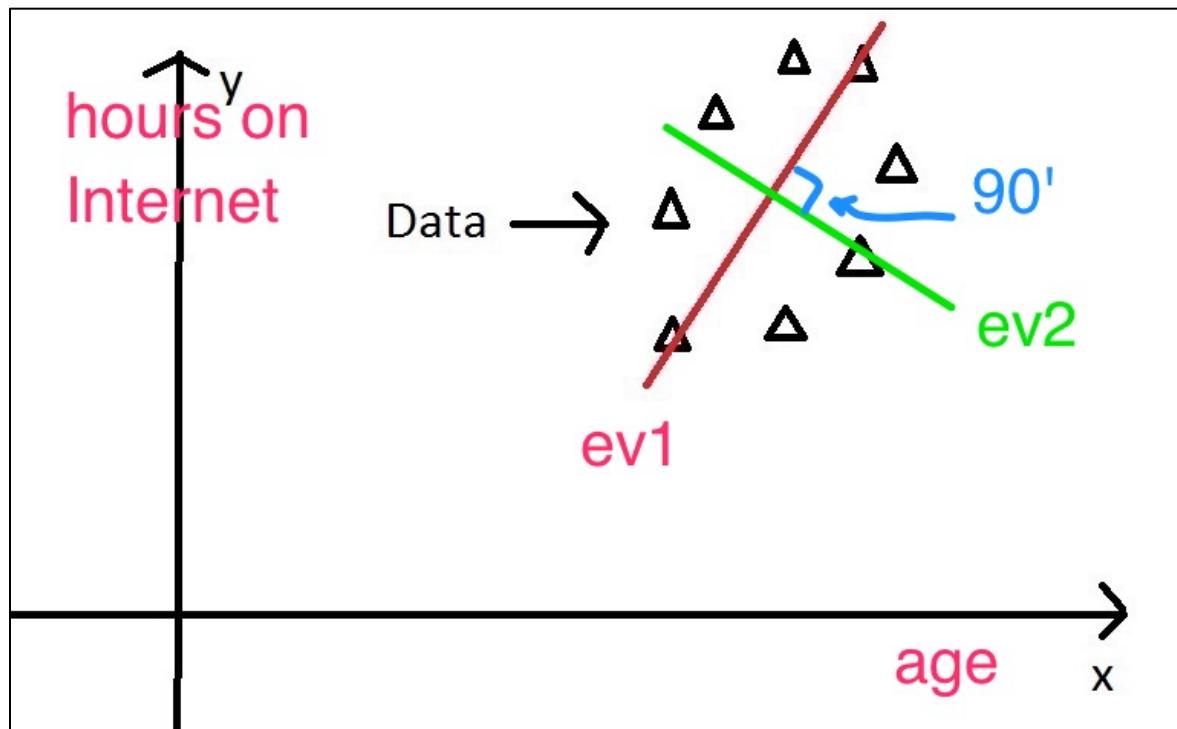
PCA Math: Eigen Vectors

- In the plot below, EigenVector (ev1) is shown that crosses the data with 'highest variance'



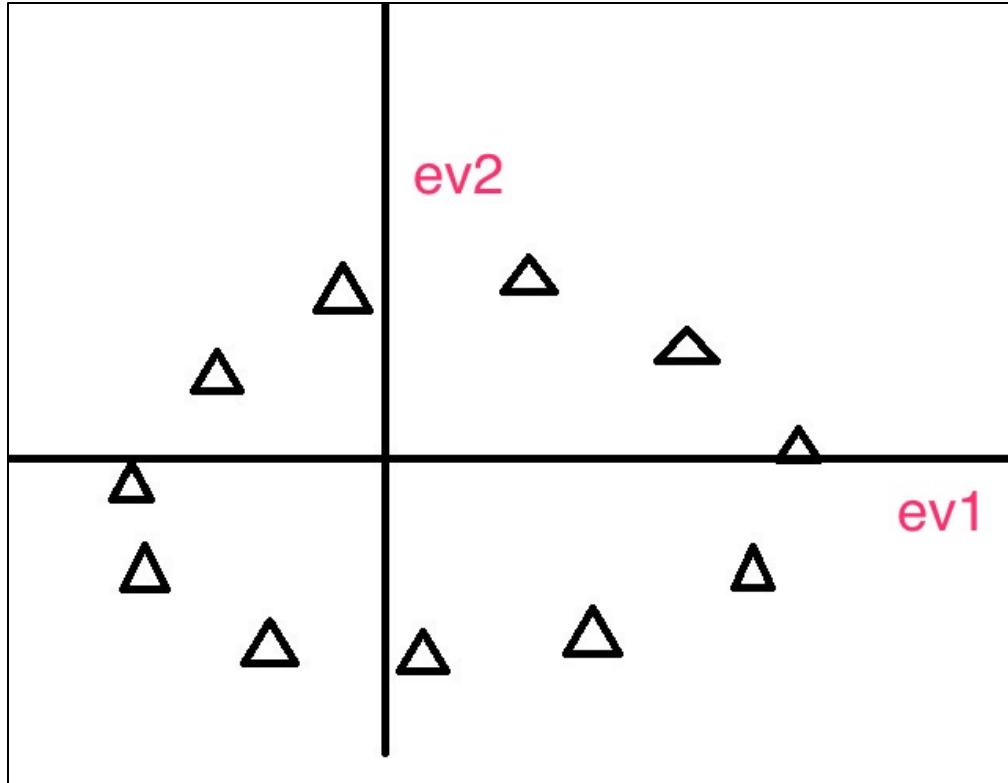
PCA Math: Eigen Vectors

- ◆ But we have a 2-dimensional data
→ 2 eigen vectors
- ◆ To maximize coverage the second EigenVector will be orthogonal (90 degrees) to the first one (ev1)



PCA Math: Eigen Vectors

- ◆ The EigenVectors have given us more useful axis to frame data
- ◆ Remember, the data hasn't changed at all, we are just looking at it from a different perspective



PCA on Oxford Internet Study

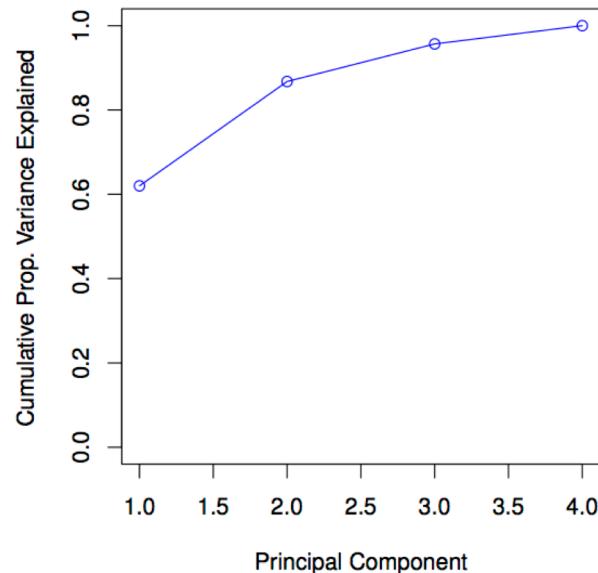
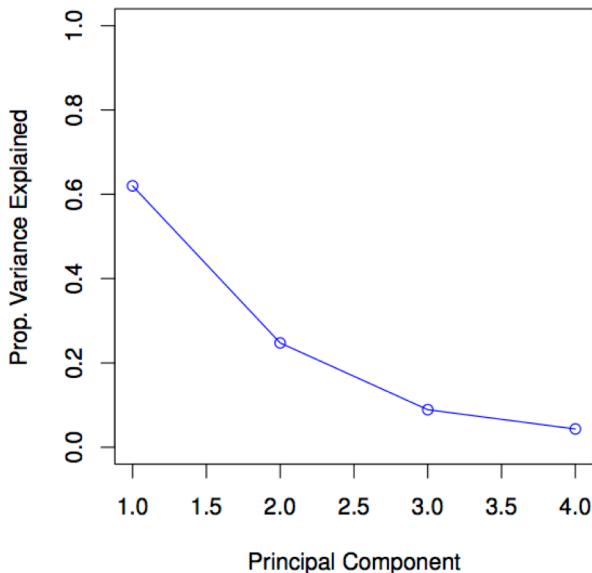
- ◆ Say there were
 - 2500 people interviewed == 2500 observations
 - 15 questions = 15 dimensions = 15 EigenVectors & Values
- ◆ Say our 15 Eigen Values are (in descending order)
[**25, 22, 20, 17**, 10, 7, 6, 5, 4, 3, 1, 0.5, 0.4, 0.3, 0.1]
- ◆ We see the first 4 have the biggest values
 - Indicating 4 directions with lots of information
- ◆ We have identified our 4 Principal Components (PC)
- ◆ We have gone from 15 attributes to 4 (call them PC1, PC2, PC3, PC4)

Principal Components

- ◆ First Principal Component (PC1) has the largest variance (EigenValue)
- ◆ Second Principal Component has second largest variance that is uncorrelated to PC1
 - Orthogonal vector to PC1

Evaluating PCA – Scree Plot

- ◆ We use a **scree plot** to understand PCA
- ◆ Left chart plots Variance for each PC component.
 - First component (PC1) has the most (62%)
 - Second one PC2 around 25%
 - PC4 has the least
- ◆ Right graph shows 'accumulated' variance when combining PCs. We are moving towards ONE

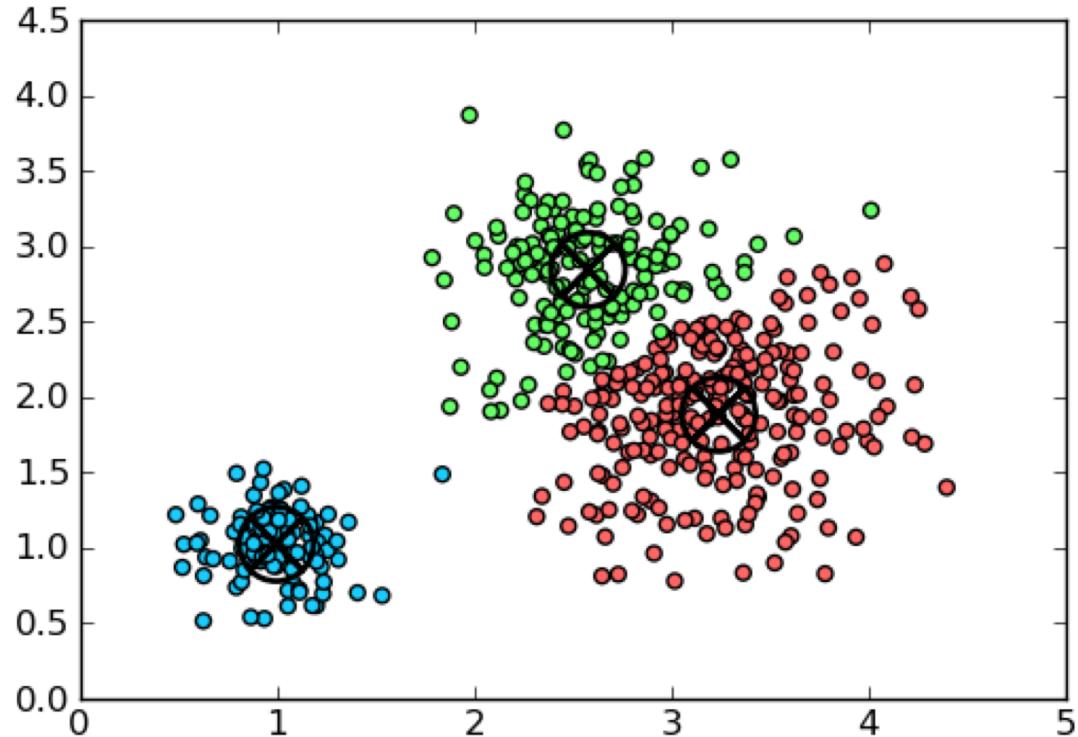


Clustering

Introduction
Supervised Machine Learning
Clustering
Recommendations
Classifications

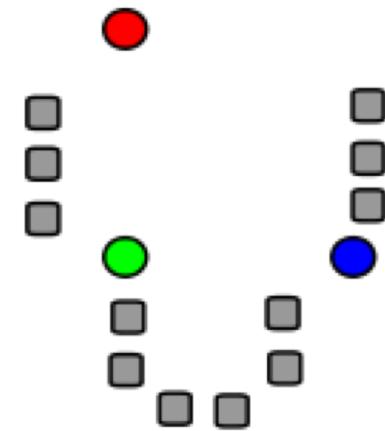
Clustering Vectors

- ◆ There are many different clustering algorithms for vectors.
- ◆ Simplest is k-means
- ◆ K-means requires a known value of k (number of clusters) to start with.

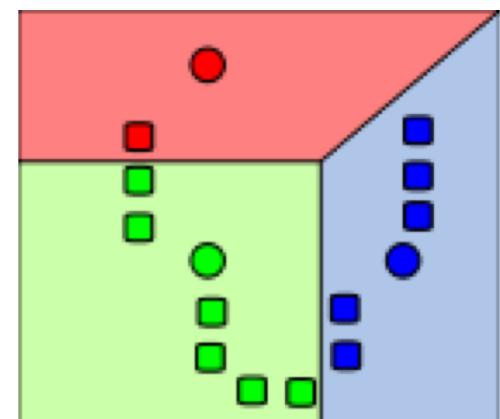


K-Means Clustering

- ◆ K-Means: Simplest Clustering Algorithm.
- ◆ Step 1: k numbers of points (centroids) are pre-seeded in the data. Example: 3 centroids (red, green, blue)

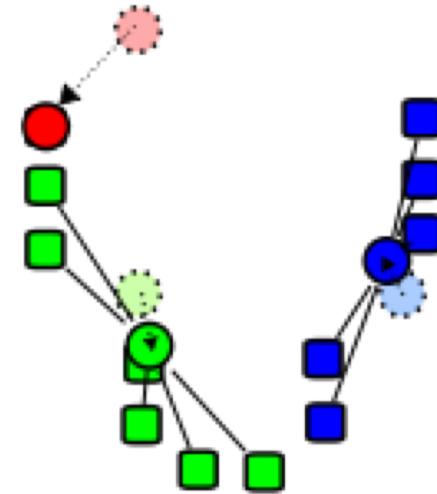


- ◆ Step 2: Each point in the dataset is associated with its nearest centroid, as determined by a distance measurement.

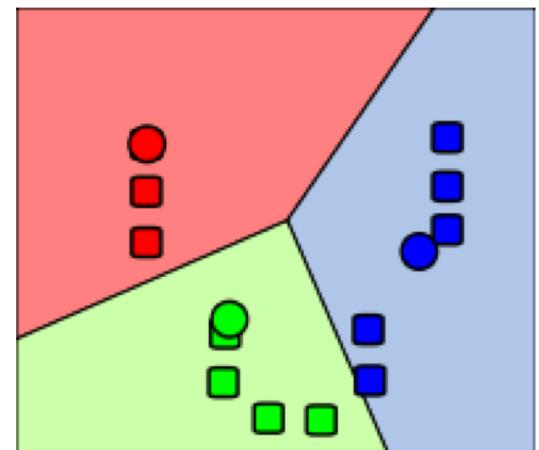


K-Means Clustering

- ◆ Step 3: The centroid (geometric center) of the clustered points becomes the new centroid of that cluster. Each centroid is updated.



- ◆ Step 4: Repeat steps 2 and 3 until convergence is reached (the points move less than the threshold amount).



K-Means Clustering Summary

- ◆ K-Means is simple

- Well-understood.
 - Easy to parallelize

- ◆ Disadvantages:

- Value of k must be known in advance, which may mean running the exercise many times to get optimum results.
 - Initial centroid positions are important; may cause long convergence.
 - Dense groupings of points are not especially considered
 - Outliers may bias results.
 - Clusters not broadly (hyper)spherical don't work well for k-means.
 - Use hierarchical clustering for these situations.

K-Means in Scikit-learn

- ◆ Scikit-learn has good support for k-means
- ◆ How to perform k-means clustering in Scikit-learn
 - Put data in ndarray
 - Perform clustering with a specified number of iterations
 - Evaluate the “fit” of the cluster. Is it a good run?
 - If not, change the number of clusters (value of k)
 - Once we have a good clustering run:
 - Map each vector to its nearest cluster
 - Group original data by its corresponding cluster
 - Assign (predict) new vectors to their nearest cluster

K-Means in Scikit-learn

- ◆ k-means clustering

```
from sklearn.cluster import kmeans
```

- ◆ Default values: just provide KMeans(df, value-of-k)

```
kmeans = KMeans(n_clusters=2,  
random_state=0).fit(X)
```

- ◆ How many iterations should you have?
 - “It depends”... too low and you might get bad results
 - Too many and you waste time
 - Try 10-20

Evaluating Cluster Performance

- ◆ WSSSE: Within Set Sum of Squared Errors

WSSSE = kmeans._inertia

- COST = sum of squared distances of points to cluster center.

- ◆ What does this mean?
 - WSSSE will decrease with increasing values of k.
 - “Law of Diminishing Returns”
 - High values of k give marginal gain.
- ◆ We can iterate across k until we get good results.

The Elbow Method

- ◆ Identify the “elbow” on the curve
- ◆ Example: What value of K to select in this case?



- ◆ How to apply the model to data?
- ◆ Use predict:
 - `kmeans.predict(Vector)`

```
int clusterId =  
model.predict(Vector.dense())
```

- ◆ Predict returns an integer
 - Cluster number, i.e., 0, 1, 2, 3... (k-1).

Handling New Data

- ◆ New Data be assigned on the existing model:
 - Make a Vector out of the new Data
 - Call KMeansModel.predict(vector), to get cluster membership as a number (integer)
- ◆ The new data will not affect the existing cluster locations.

Lab : KMeans Clustering on Flights

- ◆ **Overview:** Try clustering flights dataset
 - Use NYCFlights
- ◆ **Builds on previous labs:** None
- ◆ **Approximate time:** 20-30 min.
- ◆ **Follow:** [sklearn/kmeans/README.md](#)

Classification / Regression

Introduction

Supervised Machine Learning

Clustering

Recommendations

Classification / Regression

Classification and Regression

- ◆ Classification and Regression are both supervised ML.
- ◆ Both are to make a prediction
- ◆ Classification: predict a categorical variable
 - Could be binary (yes/no)
- ◆ Regression: predict a real number
- ◆ We "fit" the model with training data and labels
- ◆ We can evaluate the model
- ◆ Then we "predict" new data.

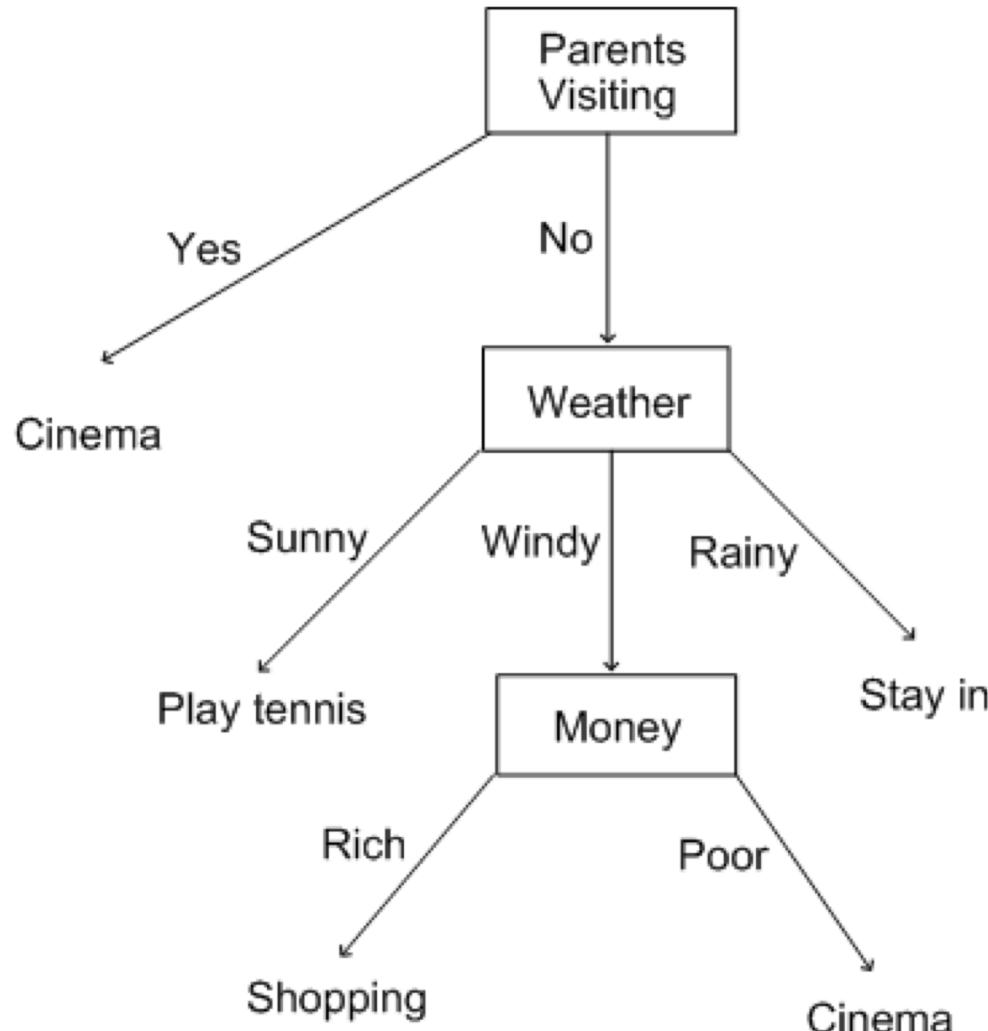
Review of Classification Algorithms

- ◆ Naïve Bayes
- ◆ Support Vector Machines (using Stochastic Gradient Descent)
- ◆ Logistic Regression/MaxEntropy
- ◆ Decision Trees
 - A basic decision tree is pretty simple.
 - It is a set of rules collect the rule
 - Decision-tree learning is a means of inferring an appropriate decision tree from the data.

Support Vector Machines (SVMs)

- ◆ Support Vector Machine is a classification method that is:
 - Supervised (trained)
 - Linear
 - Binary (splits into two classes)
(E.g., Spam or not-spam)
 - Crisp (not fuzzy, not probabilistic)
- ◆ How does it work?
 - Draw a line (hyperplane) that separates the two classes
 - Line should perform best possible separation
 - Maximize distance between line and closest points.

Decision Tree



Classification Support in Scikit-learn

- ◆ The following algorithms are well-represented in Scikit-learn:
 - Linear methods: SVM, Logistic Regression
 - Decision trees
 - Ensemble decision trees (Random Forests, Gradient Boosted Trees)
 - Naïve Bayes

Evaluating Classification Models

- ◆ Let's consider a binary classifier
 - Picks one of two outcomes (spam / not-spam)
- ◆ Two approaches
 - Confusion matrix
 - ROC curve

Confusion Matrix / Error Matrix

- ◆ Let's consider a binary classifier
 - Picks one of two outcomes (spam / not-spam)
- ◆ Say we are classifying 10 emails (6 spam, 4 not-spam)
- ◆ See '**confusion matrix**' below

Classification =>	Spam	Not Spam
Actual Spam (6 total)	4 count 66% accuracy Correct True-Positive Rate (sensitivity)	2 count 33% Incorrect False negative rate (miss rate)
Actual Not Spam (4 total)	1 count 25 % Incorrect False-positive rate (fall-out)	3 count 75% Correct True negative rate (specificity)

Confusion Matrix: More Than 2 Outcomes

		Predicted		
		Cat	Dog	Rabbit
Actual	Cat (8)	5	3	0
	Dog (6)	2	3	1
	Rabbit (13)	0	2	11

- ◆ Algorithm has trouble classifying cats / dogs (too many mis-classification)
- ◆ Can classify rabbits pretty well (less mis-classifications)

Confusion Matrix

Licensed for personal use only for Vincent Chang <vincent.chang@macys.com> from Python ML @ Macy's 04/25/2108 @
2018-04-25

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition	Positive	True positive	False negative <u>Miss rate</u> - Guilty prisoner was not convicted
	Negative	False Positive <u>Sensitivity</u> - False alarm - 'Crying wolf' - Hiring someone who is not qualified	True negative

Licensed for personal use only for Vincent Chang <vincent.chang@macys.com> from Python ML @ Macy's 04/25/2108 @
2018-04-25

Confusion Matrix : Calculations

◆ Accuracy

Overall how accurate is the model?

$$\begin{aligned} &= (TP + TN) / \text{total} \\ &= (90 + 70) / 200 \\ &= 0.8 \text{ or } 80\% \end{aligned}$$

◆ Mis-classification / Error rate

How wrong is the model?

$$\begin{aligned} &= (FP + FN) / \text{total} \\ &= (10 + 30) / 200 \\ &= 0.2 \text{ or } 20\% \\ &= 1 - \text{accuracy} \end{aligned}$$

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90)	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10)	True negative (n = 70)

Confusion Matrix : Calculations

◆ True Positive Rate (TPR)

How often model predicts 'positive' as 'positive' (**correctly**) ?

$$\begin{aligned} &= \text{TP} / \text{actual positive} \\ &= 90 / 120 \\ &= 0.75 \text{ or } 75\% \end{aligned}$$

◆ False Positive Rate (FPR)

How often model predicts 'negative' as 'positive' (**incorrectly**)

$$\begin{aligned} &= \text{FP} / \text{actual negative} \\ &= 10 / 80 \\ &= 0.125 \text{ or } 12.5\% \end{aligned}$$

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90) TPR = 75%	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10) FPR = 12.5%	True negative (n = 70)

Confusion Matrix : Calculations

◆ Specificity

How often model predicts 'negative' as negative' (correctly) ?

$$= \text{TN} / \text{actual no}$$

$$= 70 / 80$$

$$= 0.875 \text{ or } 87.5\%$$

$$= 1 - \text{FPR}$$

◆ Precision

When model predicts 'positive' how often it is right?

$$= \text{TP} / \text{predicted positive}$$

$$= 90 / (90 + 10)$$

$$= 0.9 \text{ or } 90\%$$

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90) TPR = 75%	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10) FPR = 12.5%	True negative (n = 70) Specificity = 87.5%

◆ Prevalence

How often does 'positive' occurs in our sample

$$= \text{actual positive} / \text{total}$$

$$= 120 / 200$$

$$= 0.6 \text{ or } 60\%$$

Confusion Matrix : Calculations

◆ Positive Predictive Value (PPV)

Similar to precision, takes 'prevalence' into account

$$\begin{aligned} &= \text{TP} / (\text{TP} + \text{FP}) \\ &= 90 / (90 + 10) \\ &= 0.90 \text{ or } 90\% \end{aligned}$$

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90) TPR = 75%	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10) FPR = 12.5%	True negative (n = 70) Specificity = 87.5%

◆ Null Error Rate

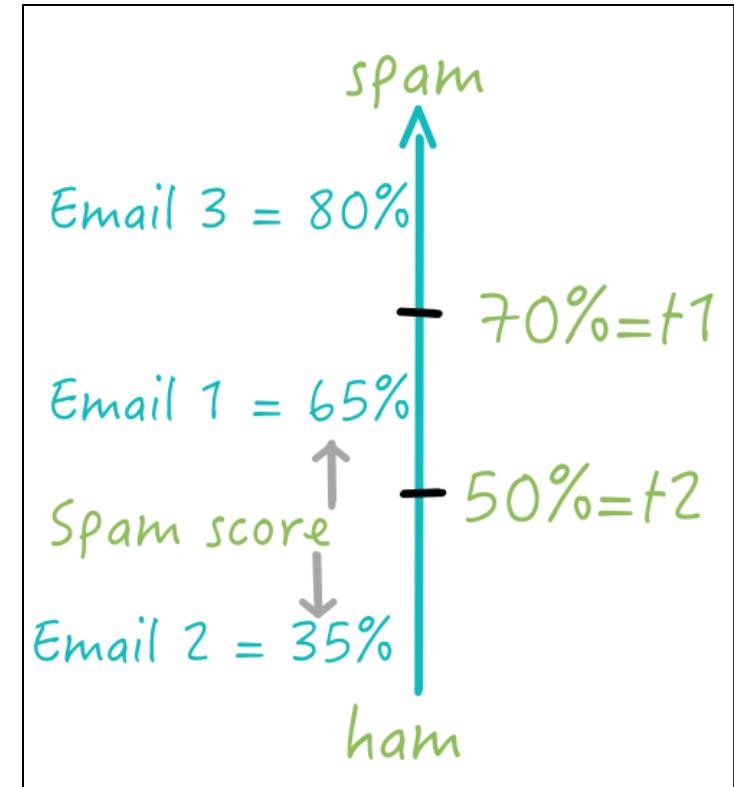
- How often the model would be wrong if it always predicted the majority class?
- Here our majority = Positive
- If we always predicted 'positive' we would be wrong 80 times (negative)
- $80/200 = 40\%$ of time

How is ROC Curve Generated

- ◆ Y-axis : True Positive Rate (TPR)
 - Actual=positive, predicted=positive
 - Correct!
- ◆ X-axis : False Positive Rate (FPR)
 - Actual=negative, predicted=positive
 - Incorrect!
- ◆ $0.0 \leq TPR \& FPR \leq 1.0$
- ◆ Plot TPR / FPR while varying 'threshold'

Threshold

- ◆ Our spam classifier provides a 'spam probability' for each email
 - Probability is between 0.0. and 1.0 (or 0 to 100%)
 - 1.0 definitely spam
 - 0.0 definitely not spam
- ◆ When an email's '**spam score**' is above a certain number we mark it as spam
- ◆ This is called '**threshold**'
- ◆ If spam threshold is lower (say 50%)
 - more emails will be classified as spam (email2, email3)
 - Users will miss emails (as they are in Spam folder)
- ◆ If spam threshold is higher (70%)
 - Fewer emails will be classified as spam (email2)
 - Users will see more spam emails be in Inbox
- ◆ We need to find the sweet spot !



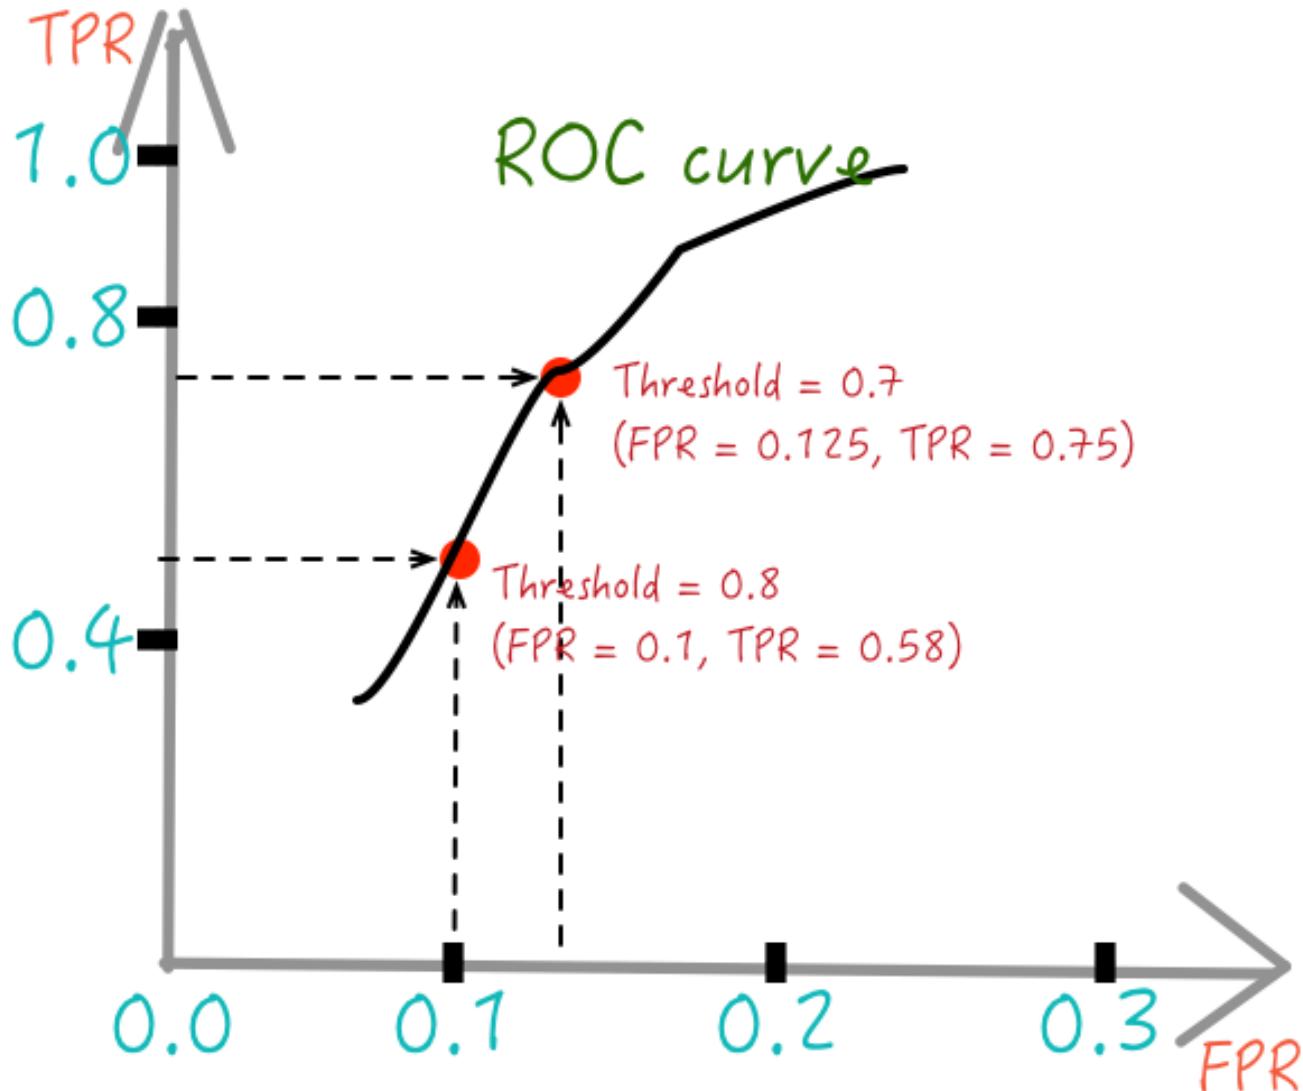
Threshold

- ◆ In first table our threshold is 0.7
 - 90 emails are correctly predicted as spam
- ◆ Next table, our threshold is higher 0.8
 - Only 70 emails are classified as spam
 - Lower TPR

		Predicted Condition	
Threshold = 0.7		Predicted Spam	Predicted Not Spam
Actual condition (total = 200)	Spam (n = 120)	True positive (n = 90) TPR = TP / positive = 90/120 = 75%	False negative (n = 30)
	Not Spam (n = 80)	False Positive (n = 10) FPR = FP / negative = 10/80 = 12.5%	True negative (n = 70)

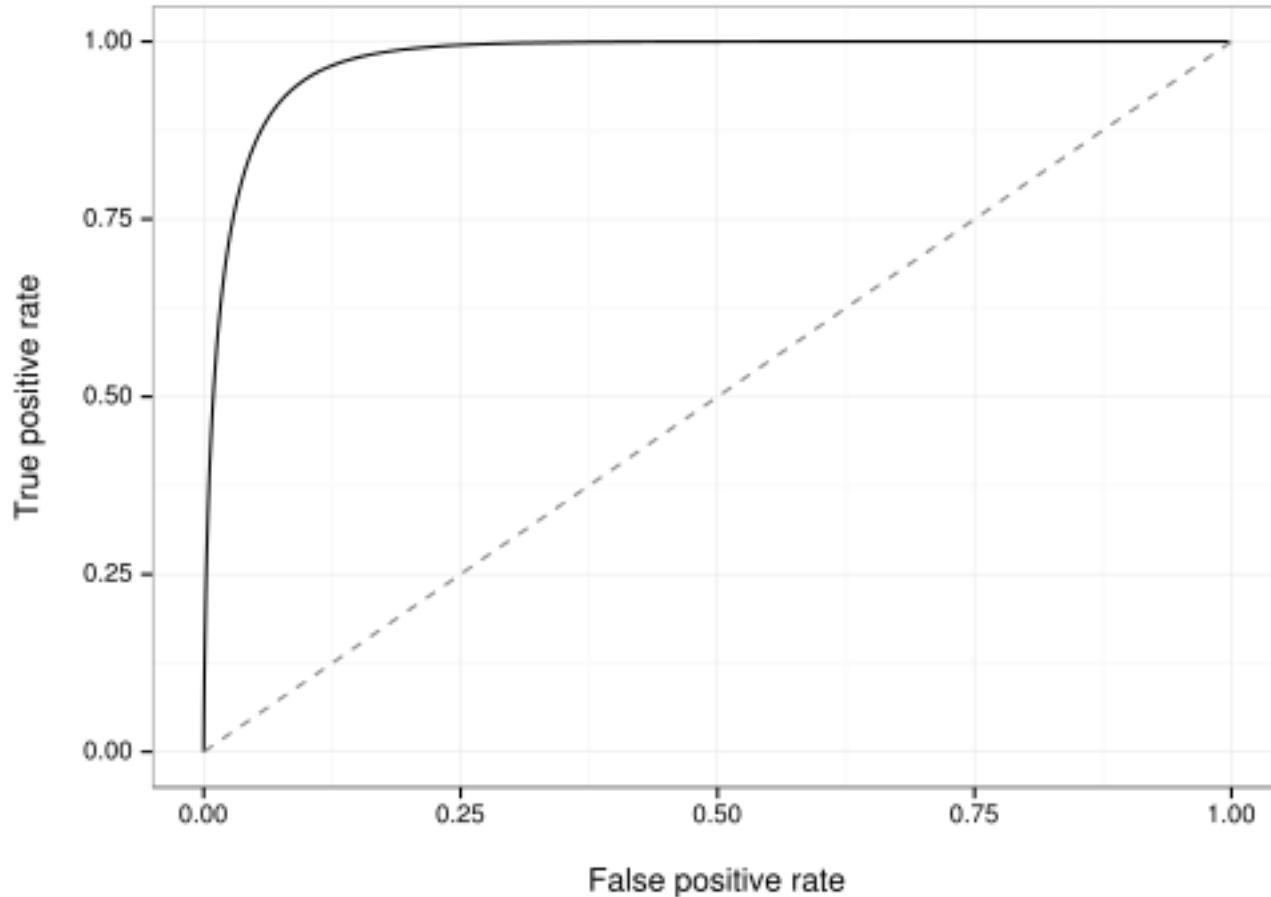
		Predicted Condition	
Threshold = 0.8 (higher)		Predicted Spam	Predicted Not Spam
Actual condition (total = 200)	Spam (n = 120)	True positive (n = 70) TPR = TP / positive = 70 / 120 = 58.33%	False negative (n = 50)
	Not Spam (n = 80)	False Positive (n = 8) FPR = FP / negative = 8 / 80 <u>= 10%</u>	True negative (n = 72)

ROC Curve : Receiver Operating Characteristic



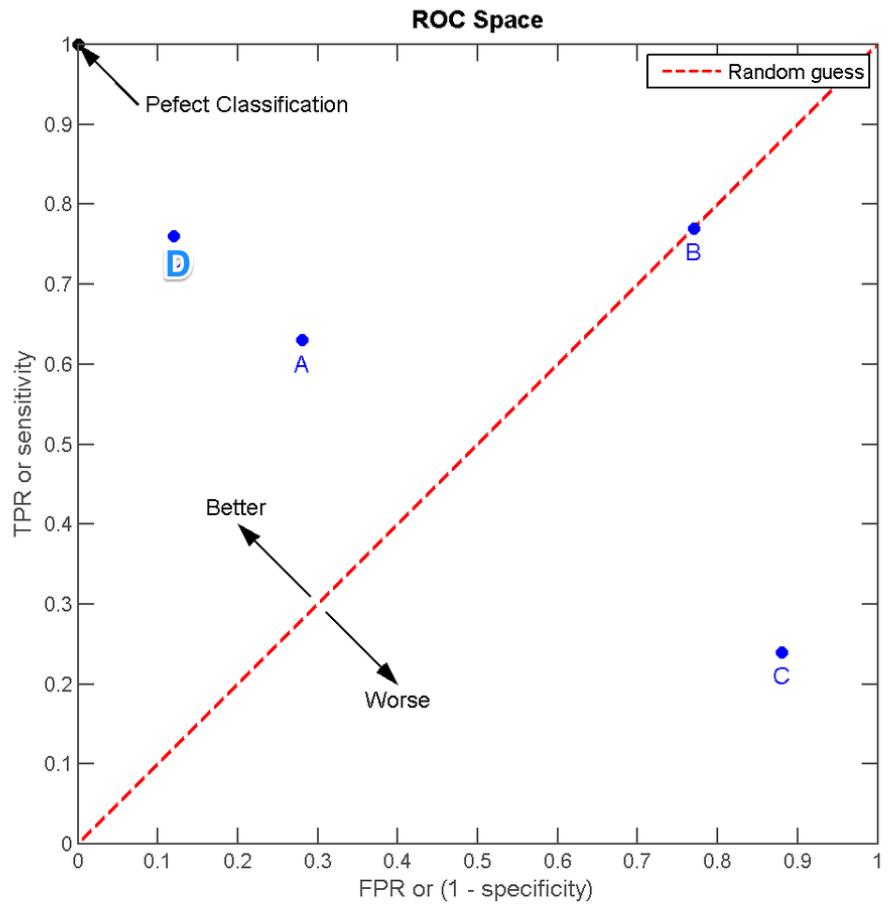
Licensed for personal use only for Vincent Chang <vincent.chang@macys.com> from Python ML @ Macy's 04/25/2108 @
2018-04-25

ROC Curve Example



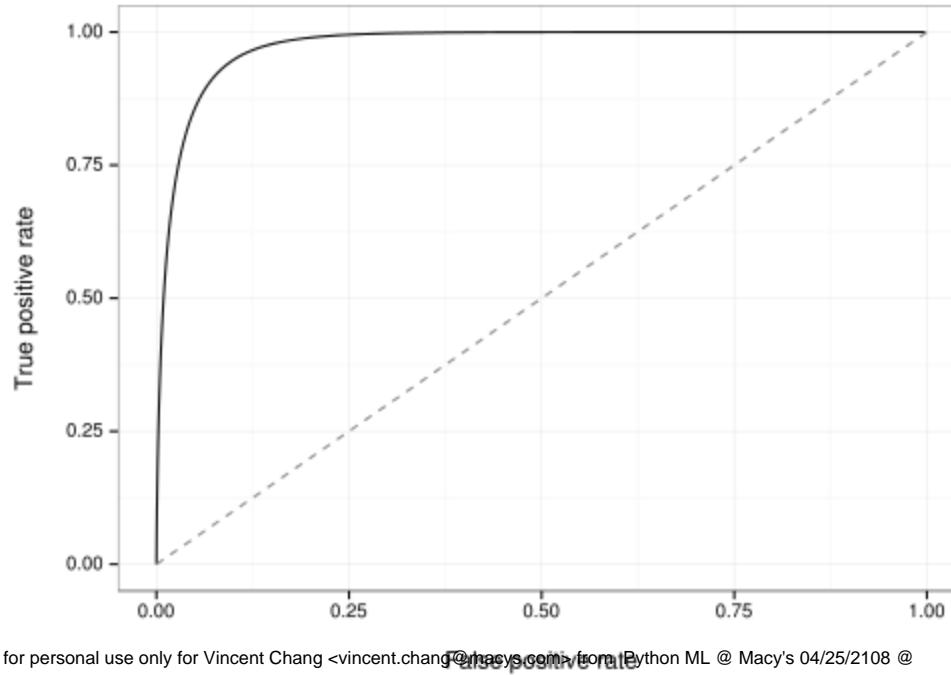
Interpreting ROC Curve

- ◆ The red-line plots 'random guess' = B
- ◆ Approaching 'top left' corner would be a perfect classifier!
So D is better A
- ◆ C performs worse than random → bad



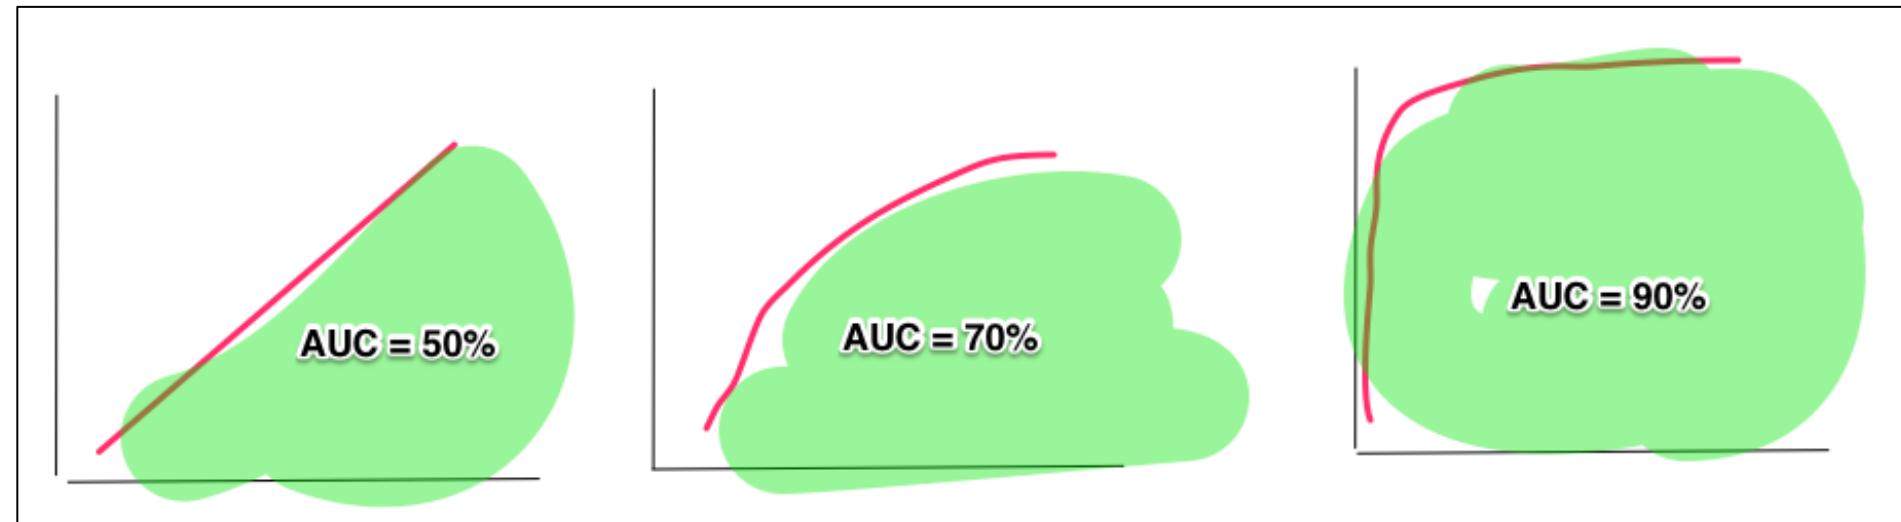
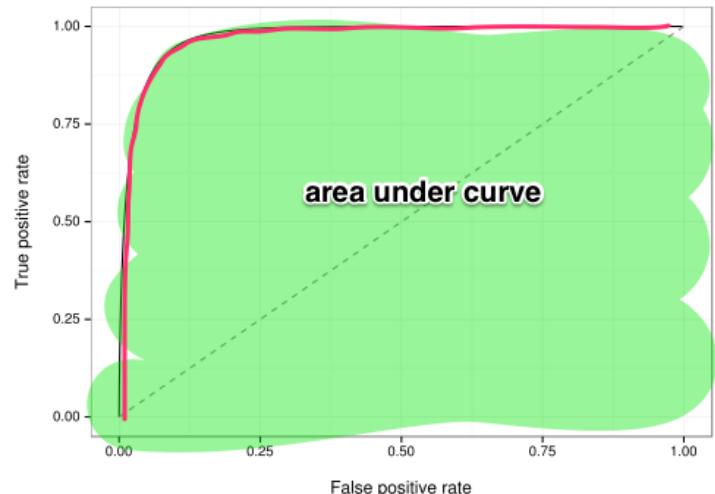
Interpreting ROC Curve

- ◆ Shows tradeoff of TPR (sensitivity) vs. FPR (1 – specificity)
- ◆ The closer to top-left , the more accurate the model
- ◆ Upper left corner (0,1) = perfect classification!
- ◆ The closer to middle line (45 degree) the less accurate the test
 - Middle line represents : random classification (50%)



Area Under Curve – AUC (ROC Space)

- ◆ Measures the percentage of area 'under the curve'
- ◆ AUC is between 0 and 1
- ◆ Higher AUC → more accurate the model
- ◆ See 3 scenarios below
 - Left most is bad (50%)
 - Middle : OK (70%)
 - Right most : very good (90%)



Using AUC to Measure Accuracy

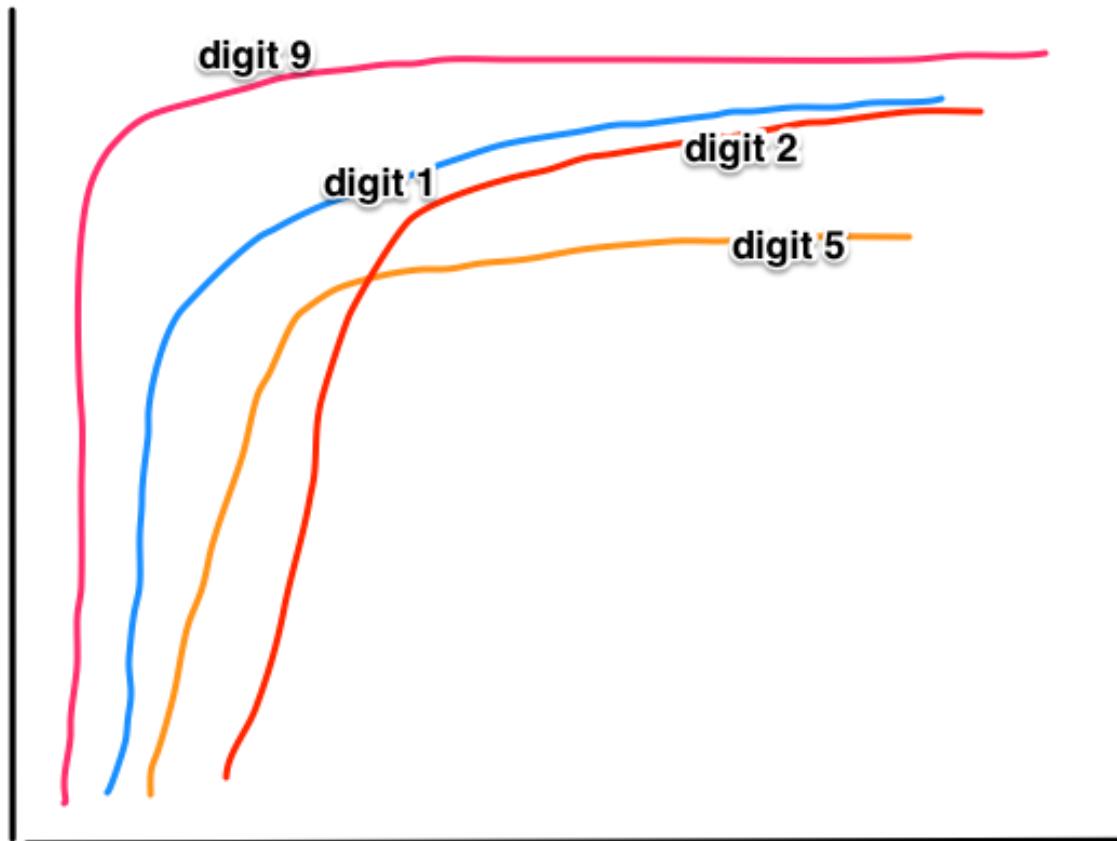
- ◆ Accuracy can be specified using a grading system

AUC	Grade
0.9 – 1.00	A - Excellent
0.80 – 0.90	B - good
0.70 - 0.80	C - fair
0.60 – 0.70	D - poor
0.50 – 0.60	F - Fail



ROC / AUC For Multiclass Classifiers

- ◆ Say our algorithm recognizes hand-written digits (postal code) into numbers.
- ◆ Its ROC can be drawn as follows



Lab: Classification

- ◆ **Overview:** In this lab, we will use the “nycflights13” dataset for predicting delayed flights
 - We will use logistic regression as a
- ◆ **Builds on previous labs:** None
- ◆ **Approximate time:** 25-35 min.

Recommendations

Introduction
Supervised Machine Learning
Clustering
Classification/Regression
Recommendations

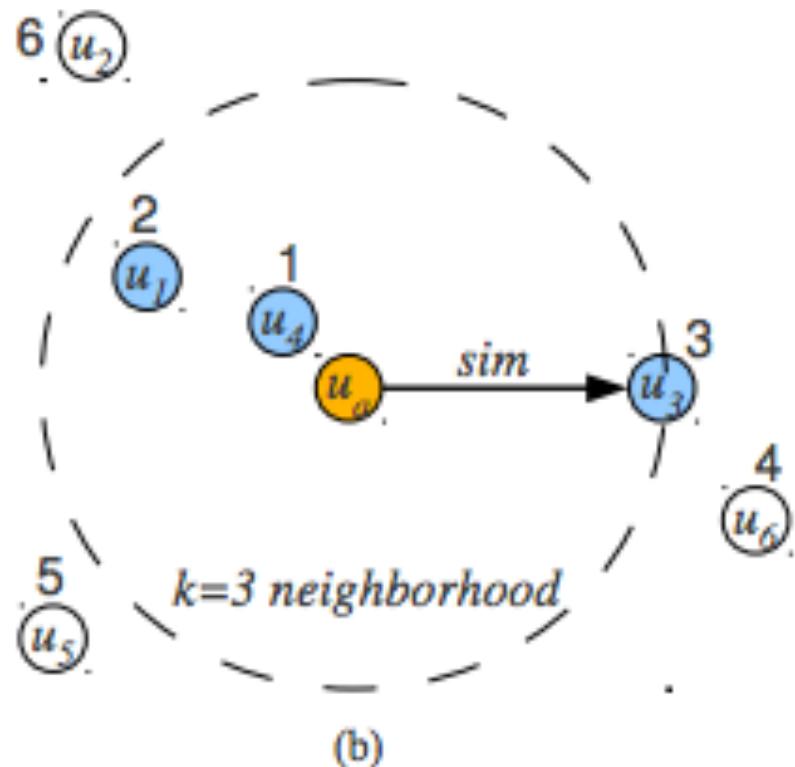
Collaborative Filtering

- ◆ Collaborative Filtering (CF) is commonly used in “Recommended For You” or “More Like This” functionality
- ◆ Recommendations can be explicit (based on user ratings), or implicit (based on user interest)
- ◆ CF is expressed as Users -> Items
 - However, any correlation could be modeled as users to items
- ◆ Users and Items could be the same (example: dating site)

Ratings Matrix: Users/Movies

	i_1	i_2	i_3	i_4	i_5	i_6	i_7	i_8
u_1	?	4.0	4.0	2.0	1.0	2.0	?	?
u_2	3.0	?	?	?	5.0	1.0	?	?
u_3	3.0	?	?	3.0	2.0	2.0	?	3.0
u_4	4.0	?	?	2.0	1.0	1.0	2.0	4.0
u_5	1.0	1.0	?	?	?	?	?	1.0
u_6	?	1.0	?	?	1.0	1.0	?	1.0
u_a	?	?	4.0	3.0	?	1.0	?	5.0
r_a	3.5	4.0		1.3		2.0		

(a)



Item Ratings Matrix - Normalized

S	i_1	i_2	i_3	i_4	i_5	i_6	i_7	i_8	u_a
i_1	-	0.1	0	0.3	0.2	0.4	0	0.1	2
i_2	0.1	-	0.8	0.9	0	0.2	0.1	0	?
i_3	0	0.8	-	0	0.4	0.1	0.3	0.5	?
i_4	0.3	0.9	0	-	0	0.3	0	0.1	?
i_5	0.2	0	0.7	0	-	0.2	0.1	0	4
i_6	0.4	0.2	0.1	0.3	0.1	-	0	0.1	?
i_7	0	0.1	0.3	0	0	0	-	0	?
i_8	0.1	0	0.9	0.1	0	0.1	0	-	5
r_a	-	0.0	4.6	2.8	-	2.7	0.0	-	

Recommendations in Python

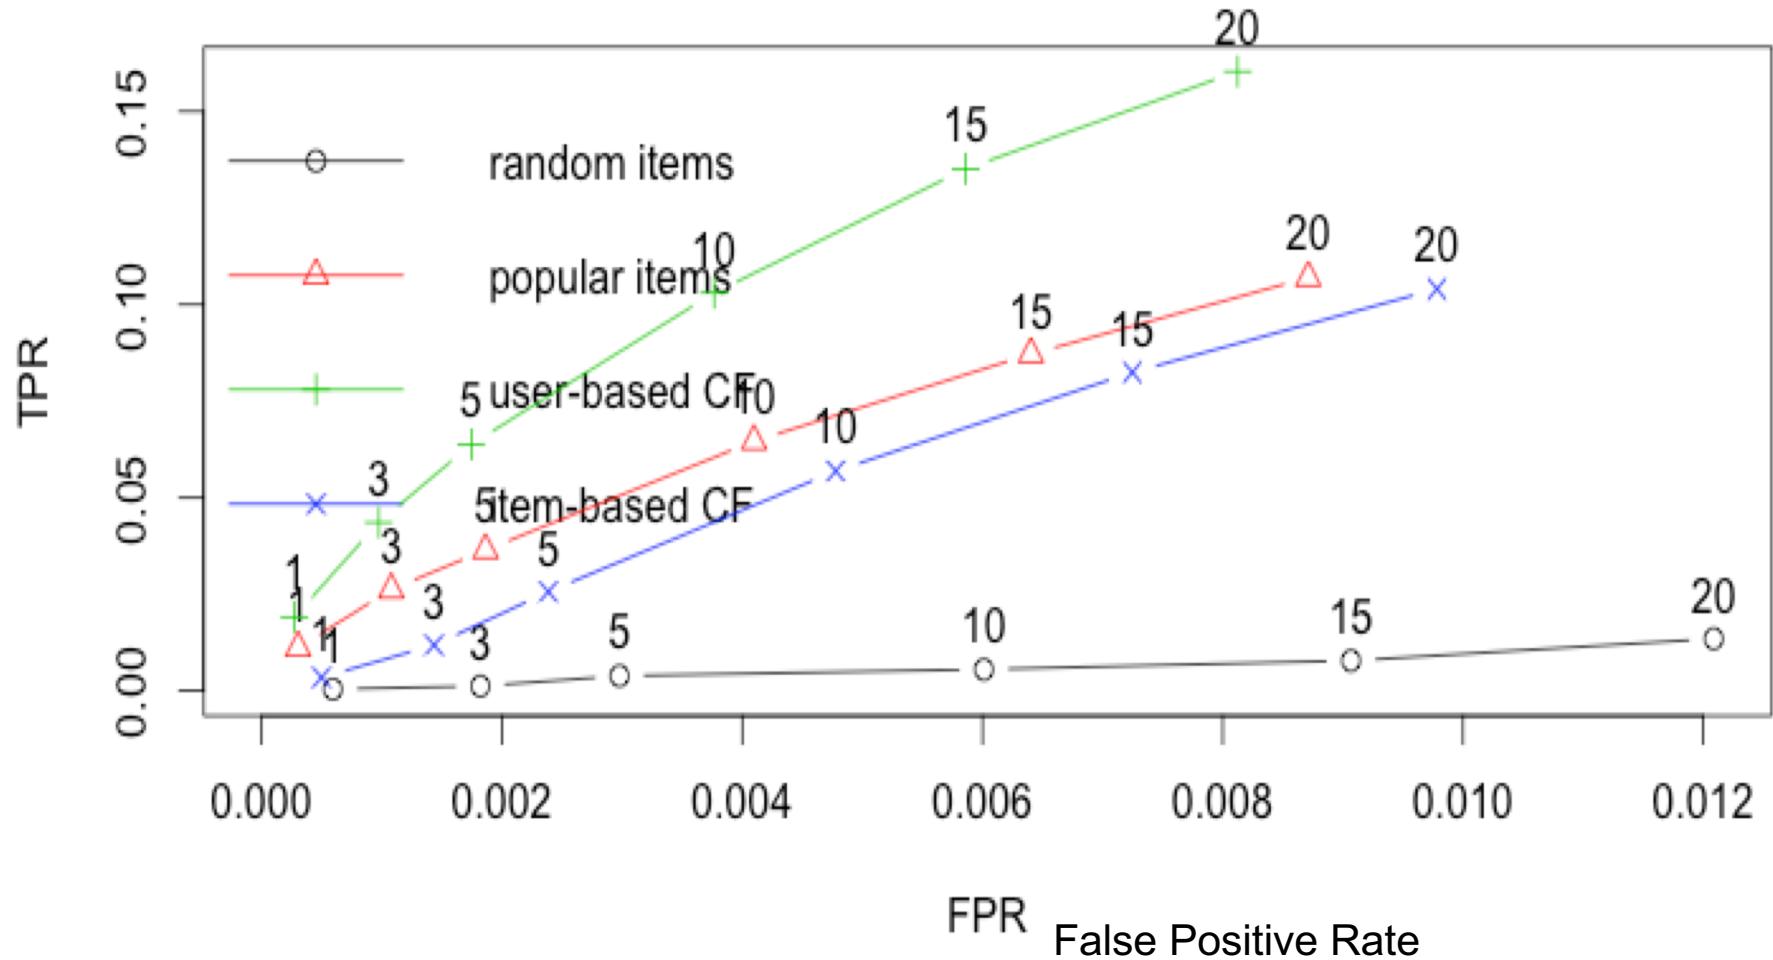
- ◆ There is a specialized library called **scikit-surprise** for Recommendations
- ◆ Has several ways of performing recommendations
 - SVD (Singular Value Decomposition)
 - ALS (approximation of SVD)
 - K-Nearest Neighbors

Preparing Data for Scikit-surprise

- ◆ Scikit-surprise only looks at integer userid and itemid.
 - Assign your data with unique integer userid and itemid.
- ◆ Preferences are expressed as a double (higher is better)
 - What if your preferences are binary (yes/no)?
 - Assign a number, say, 5.0 for yes, and 1.0 for no.
 - What if preferences are just “implied”?
 - User viewed item as 5.0, no data for unviewed item.
- ◆ Scikit-learn Rating: Integer, Integer, Double

Evaluating CF Recommendations (ROC)

True Positive Rate



Lab : Recommendations

- ◆ **Overview:** We will use MovieLens example
 - There are a number of TODOs to finish in the code
- ◆ **Builds on previous labs:** None

- ◆ **Approximate time:**
20-30 min.
- ◆ Follow:
6-Scikit-learn/recs/README.md